

Vishwas Kumar

Software Engineer

956-928-8612 ◇ vishwas101kumar@gmail.com ◇ Hyderabad, Telangana ◇ [LinkedIn](#)

SUMMARY

Java Backend Developer with 2 years experience building REST APIs and microservices. Skilled in Spring Boot, AWS, and Kafka. Focused on clean, testable code and practical caching optimizations to deliver reliable solutions.

SKILLS

Languages & Concepts Java, C++, JavaScript, DSA, Agentic AI

Frameworks & Libraries: Spring Boot, Hibernate/JPA, JUnit/Mockito, ReactJS, Redux

Data & Streaming: MySQL, PostgreSQL, Redis, Apache Kafka, Apache Spark

Cloud, DevOps & Tools AWS, Docker, Kubernetes, Terraform, GitLab CI/CD, Git, Datadog, Grafana, K6

EXPERIENCE

Associate Software Engineer

Demandbase

Aug '25 — Present
HYDERABAD, India

Project: Match Mapping Near Real Time (NRT) Pipeline

- Designed and developed a **Near Real Time (NRT)** data pipeline using **CDC Architecture** from **PostgreSQL** to **CelerData (StarRocks)**, reducing data inconsistency time from **24 hours to less than 5 seconds** for different customer tenants.
- Developed a high-performance **Spark** job as part of the NRT pipeline to process **240+ million** records from **Kafka** topics and efficiently ingest them into **Apache Iceberg** tables.
- Developed a **Kafka Connect** application using **Debezium** to capture PostgreSQL WAL (Write-Ahead-Log) changes from 10+ tables and stream them to Kafka topics for real-time data processing.

Project: Account Intelligence Agent (AIA)

- Built a multi-tenant **AI agent** (FastAPI, LangGraph, AWS Bedrock / Claude) with 7+ tools and streaming APIs for an AI sales assistant.
- Built a 5-stage **AI guardrail pipeline** (PII, and intent filtering) to block prompt-injection and adversarial inputs.
- Added multi-turn memory on **AWS AgentCore** with a circuit breaker and **selective pruning**, preventing conversation loss during outages.
- Led a zero-downtime **LiteLLM proxy** migration behind LaunchDarkly flags, enabling safe per-tenant rollout with instant rollback.

Project: Sales-AI-Summary

- Integrated LLM providers (**AWS Bedrock**, **LiteLLM**) with LaunchDarkly flags for dynamic model routing, content-safety guardrails, and zero-downtime rollouts.
- Built **REST APIs** on PostgreSQL/Spring Data JPA with 22 Liquibase migrations and parameterized queries to prevent SQL injection.
- Designed an async pipeline (**Redis** → **AWS SQS** → multi-threaded consumer) to refresh AI summary caches at scale, using ShedLock to avoid overlapping runs across replicas.
- Built a **Strategy-pattern** summary engine with 7 pluggable topic providers and tiered Redis TTL caching to cut redundant LLM calls.

SDE Intern

Demandbase

Jul '24 — Jul '25
HYDERABAD, India

- Resolved a critical CRM data duplication bug** by refactoring the data transformation pipeline to enforce strict data mapping and integrity.
- Modernized cloud deployments** by migrating 3 microservices to **Helm charts** and configuring **Kubernetes** Pod Disruption Budgets to ensure high availability.
- Engineered API performance tests** using **Grafana K6** to simulate high-load bulk operations and identify backend scalability bottlenecks.
- Enhanced system reliability** by writing comprehensive unit tests (**JUnit/Mockito**) and integrating automated API validations into the **GitLab CI/CD** pipeline.

EDUCATION

Bachelor of Technology in Computer Science and Engineering, Lovely Professional University (GPA: 8.12/10.0)

Jul '21 — May '25
Phagwara, Punjab, India